

# AN10775

## NicheLite for LPC implementation notes

Rev. 01 — 19 December 2008

Application note

### Document information

Info	Content
<b>Keywords</b>	Network, Ethernet, TCP/IP Stack, LPC2400, LPC3250
<b>Abstract</b>	This application note discusses implementation details when using the NicheLite for LPC TCP/IP stack in a project. Aspects discussed include memory management, stack operation and customization.

**Revision history**

Rev	Date	Description
01	20081219	First version.

**Contact information**

For additional information, please visit: <http://www.nxp.com>

For sales office addresses, please send an email to: [salesaddresses@nxp.com](mailto:salesaddresses@nxp.com)

1. Introduction

Several members of the NXP LPC2000 series flash microcontrollers provide an on-chip Ethernet controller, allowing easy connection to a network. NXP makes available for LPC2000 users the NicheLite for LPC TCP/IP stack. This stack is a cut down version of the full stack available from InterNiche.

The advantage of using the NicheLite for LPC stack is that it is royalty free and it is already ported to the LPC2000 family. However embedded engineers still have to integrate the stack with their application. Integration considerations include performance and memory usage.

This application note assumes basic knowledge of TCP/IP protocols, such as IP, TCP and UDP. It is intended to complement the documentation from InterNiche, and not replace any such documentation.

2. Packets

This section describes aspects related to packets. Data is transferred across the network in packets, which can be up to the Maximum Transfer Unit (MTU) in size.

2.1 Packet Allocation and Deallocation

All data in and out of the TCP/IP stack is transferred using packets. Packets contain the Ethernet header, the IP header and additional data depending on the protocols used. The packets are allocated from the heap when the stack starts executing and remain allocated for the runtime of the firmware. All packets are placed into a pool of free packets after allocation. They are moved out of the free pool when they are used by the stack and back into the free pool when they are no longer needed by the stack.

2.1.1 Packet Types and Configuration

There are two types of packet in the NicheLite implementation, big and little ("lil" in the stack source code). The number and maximum size of each type is defined in ipport.h.

Table 1. Packet Configuration Identifiers

Identifier	Description
NUMBIGBUFS	Total number of big packets
BIGBUFSIZE	Size of big packets in octets
NUMLILBUFS	Total number of little packets
LILBUFSIZE	Size of little packets in octets

When a packet is needed by a protocol pk\_alloc is called. This function uses a little packet if the data size for the packet is small enough, otherwise it uses a big packet. If all the little packets are in use then a big packet will be used.

The identifiers can be adjusted to optimize packet usage for the specific embedded system. If a lot of small packets will be sent then it may be beneficial to increase the number of little packets and reduce the number of big packets, for example.

## 2.2 Increasing the Total Number of Packets

The maximum packet size (MTU) is set to a typical value, such as 1514 bytes. This allows data up to the MTU in size be transmitted across most networks without fragmentation.

It can be desirable to increase the number of packets available in an embedded system to allow for increased concurrent communications to take place. Care must be taken to ensure that at no point in time all packets are in use and another is needed.

For example an embedded system may have a TCP based control channel, an IGMP management system and a TELNET based interface. In the worst case all three communication interfaces will be transmitting and receiving packets at the same time. The problem becomes more critical if the embedded system needs to stream unconfirmed packets of data.

The solution is to increase the number of packets available to the stack. This is possible by decreasing the maximum packet size that the embedded system can handle.

By changing the MTU the system will not be able to transmit or receive packets above that size, but more packets can fit into the available memory. Configuring the maximum packet size in the NicheLite stack configures the Ethernet Maximum Frame Register (MAXF), which will cause the Ethernet controller to reject packets above this size without interrupting the CPU.

How can the maximum packet size be determined? This is achieved by examining the existing protocols that will be needed, such as DHCP, along with careful design of custom protocols. The following table indicates the typical data sizes for commonly used protocols.

**Table 2. Typical Protocol Data Sizes**

Protocol	Data Size (octets)
DHCP (must have a minimum of 312 octets of options)	548
ARP	28

To these values the UDP header, IP header and Ethernet header lengths must be added to get the largest packet size. Some examples are in the following table.

**Table 3. Typical Header Sizes**

Protocol	Header Size (octets)
UDP	8
IP	24 (depending on options used)
Ethernet	22

For example DHCP packets might be  $548 + 8 + 24 + 22 = 602$  octets in size. Note that not all protocols are built upon UDP.

### 2.2.1 Stack Configuration

To change the maximum size of packets supported by the Ethernet controller and the stack:

- In `ipport.h` change `BIGBUFSIZE` to the maximum packet size, including all headers

- In ether.h change MTU and ET\_MAXLEN to BIGBUFSIZE minus 22 (the size of the Ethernet header)

## 2.2.2 Testing

It is recommended to stress test the embedded system in order to produce the worst case scenario in terms of packets in use at any moment in time.

## 2.3 Runtime Monitoring

The usage of packets can be monitored during runtime when using a suitable debug interface, such as JTAG or an In-Circuit Emulator.

Access to some internal data is also available via the default serial interface system, however because the interface may affect performance and may not be needed or desired for production firmware, it may be disabled. Using a system such as JTAG provides more flexibility during debugging, and is recommended.

### 2.3.1 Detailed Packet Usage

In ipport.h ensure NDEBUG is defined. This allows access to the pktlog array, defined in pktalloc.c. Create a watchpoint on pktlog to access the contents when execution is stopped.

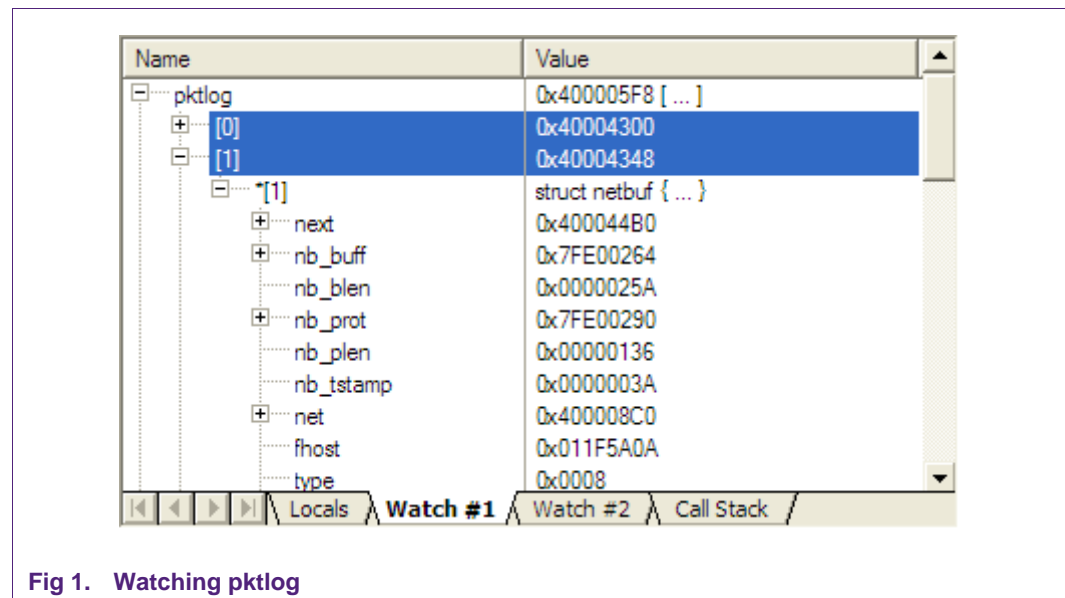


Fig 1. Watching pktlog

The first NUMBIGBUFS (see ipport.h) entries in the array are the big packets. The remaining entries are the little packets. For example if NUMBIGBUFS is set to 10, then packets zero to nine will be big, followed by the little packets.

The packet data is located at nb\_buff. This includes Ethernet and IP headers. Higher layer protocol data is stored at m\_data, and has a length of m\_len. If a packet is currently in use by the stack then the inuse member will be non-zero. The fhost member contains the IP address of the remote host.

By looking at fhost, inuse, m\_data and m\_len it is possible to understand the purpose of all the packets in the system when execution is paused.

All members of the netbuf structure are briefly described in netbuf.h.

### 2.3.2 Total Packet Count

By adding a few lines of code to pktalloc.c it is possible to make debugging packet usage a little easier. The method involves creating two variables to hold the current number of packets of each type in use.

In pktalloc.c create two variables to keep count.

```
1  int bigbuffree = NUMBIGBUFS;
2  int lilbuffree = NUMLILBUFS;
```

In the pkt\_alloc function add two lines to decrement the counters.

```
3  if ((len > lilbufsiz) || (lilfreeq.q_len == 0)) /* must use a big buffer */
4  {
5      p = (PACKET)getq(&bigfreeq);
6      if (p) bigbuffree--;
7  }
8  else
9  {
10     p = (PACKET)getq(&lilfreeq);
11     if (p) lilbuffree--;
12 }
```

In the pkt\_free function add two lines to increment the counters.

```
13 if (pkt->nb_blen == bigbufsiz)
14 {
15     q_add(&bigfreeq, (qp)pkt);
16     bigbuffree++;
17 }
18 else
19 {
20     q_add(&lilfreeq, (qp)pkt);
21     lilbuffree++;
22 }
```

The counters can be watched during runtime to monitor the usage of both packet types under various conditions.

### 2.3.3 Packet Statistics

The NicheLite stack provides a structure containing statistics on the packets. This includes such aspects as the number of packets transmitted and received and the number of packets discarded. This is very useful for determining if there might be unseen data processing problems.

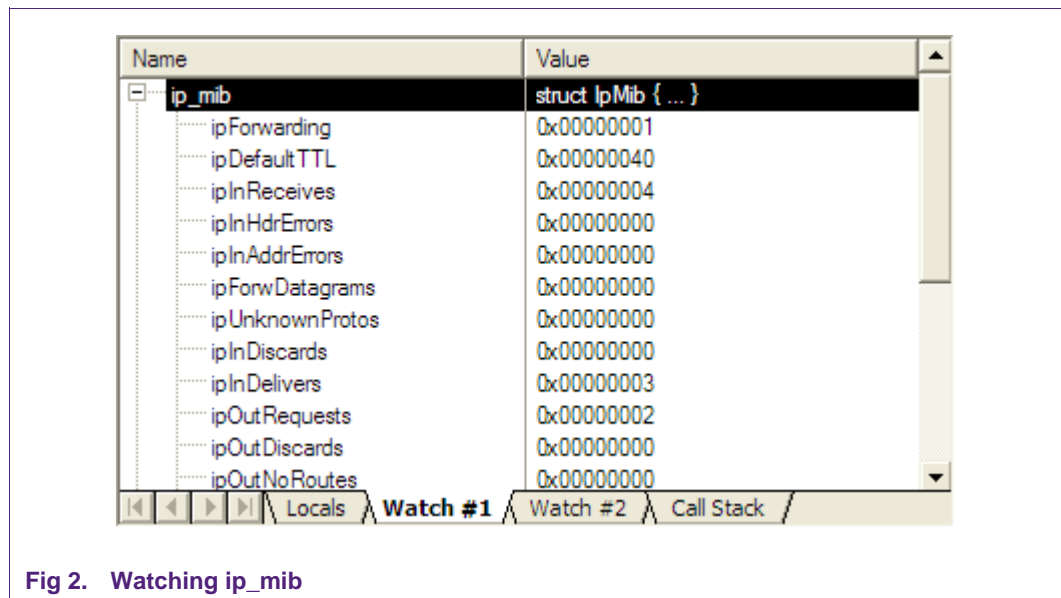


Fig 2. Watching ip\_mib

The ip\_mib structure is defined in m\_ip.c.

### 3. TCP

This section describes various aspects of the TCP implementation in the NicheLite stack and how to accomplish specific tasks.

#### 3.1 Retransmits

If the embedded system periodically sends segments of data to a remote host, and the segments are not acknowledged, then they are automatically retransmitted up to 12 times. The time between retransmits increases each time. During this time the memory remains allocated for the Ethernet packet containing the segment.

TCP has the ability to transmit additional segments while handling retransmits for other segments. If the connection is broken or the remote host crashes, there could be an increasing backlog of segments going through the retransmit process. In the worst case all Ethernet packet buffers will be exhausted. This problem could be exacerbated if the embedded system periodically transmits TCP segments.

It is possible to change the number of retransmits using the TCP\_MAXRXTSHIFT identifier defined in mtcp.h. Whether this is a suitable change depends on the application being developed.

It is recommended that a TCP based protocol implement a command/response sequence for controlling the flow of data. This would allow failure of the network or remote host to be detected before additional TCP segments are transmitted.

Note that closing a connection causes the stack to free all pending packets for the connection.

An alternative method is to use TCP keepalives to detect a broken connection.

3.2 Keepalives

Keepalives are small messages periodically sent between two hosts to inform each end that the connection is operational.

Suppose a remote host opens a TCP connection to the embedded system. While the connection is open there is no indication that the connection is functional unless data is sent. If the remote host crashes or the network connection is broken the embedded system will not know about it. This can cause two potential problems.

- Any packets in the process of being assembled for transmission to the remote host can get stuck in limbo, causing a memory leak
- The socket for the connection to the remote host is never closed, causing a memory leak

Although the memory leak may be slow, if the embedded system is required to have 100% uptime, the system could run out of RAM after months of operation.

There are three aspects to using TCP keepalives in the NicheLite stack.

- How to enable the feature for specific sockets
- The default timing follows RFC1122, which is probably too long for most embedded systems
- There is a minor issue in the implementation in some versions of the stack

The following subsections deal with each of these aspects in turn.

3.2.1 Enabling Keepalives

Typically socket options are enabled and disabled using the `m_ioctl` function, however the NicheLite implementation of this function does not support enabling keepalives. Therefore the feature must be enabled directly. Alternatively the `m_ioctl` function is easily modified to add this option.

Calling `m_listen` or `m_connect` returns a pointer to a socket. The option is enabled by using the `so_options` member.

```
23 WP_SOCKETTYPE sock;
24 sock = m_listen(&sin, upcall, &e);
25 // enable keep alives
26 sock->so_options |= SO_KEEPAIVE;
```

3.2.2 Adjusting Timing

As previously mentioned, the default timing for keepalives follows the Internet standard RFC1122. The timing values are defined in `mtcp.h` and are listed in the following table.

Table 4. TCP Keepalive Timing Identifiers

Identifier	Description
TCPTV_KEEP_INIT	The time during which the connection must be established
TCPTV_KEEP_IDLE	Default time before probing
TCPTV_KEEPINTVL	Default probe interval



Identifier	Description
TCPTV_KEEPCNT	Max probes before drop

When modifying these values be aware that TCPTV\_KEEP\_INIT may be defined twice, depending on the version of the NicheLite stack.

As an example, the default value of TCPTV\_KEEP\_IDLE is two hours. This means that if no data is received for two hours then keepalives will be transmitted. In a typical embedded system a more appropriate value might be one minute.

The number of keepalive probes that go unacknowledged before the stack drops the connection is also configurable. The best strategy is probably to work out the maximum time that the system can commit RAM to a dead connection, and then work out appropriate keepalive values to achieve that requirement.

Note that these timing settings apply to all connections that have keepalives enabled.

### 3.2.3 Fixing Sequence Number Issue

This section might not apply to your version of the NicheLite stack.

RFC1122 states that the keepalive message must have a sequence number of one less than the next expected sequence number. The NicheLite stack may send the next sequence number instead. Those keepalives will not be acknowledged by some TCP/IP stacks, including the implementation in Windows XP.

In function tcp\_output in tcpout.c, change the following lines

```

27  /* set seq for a BSD-ish keepalive */
28  ptcp->th_seq = ((tp->snd_nxt >> 24) & 0x00ff) |
29                ((tp->snd_nxt >> 8) & 0xff00) |
30                ((tp->snd_nxt & 0xff00) << 8) |
31                ((tp->snd_nxt & 0x00ff) << 24);

```

to the following

```

32  /* set seq for a BSD-ish keepalive */
33  ptcp->th_seq = (((tp->snd_nxt - 1) >> 24) & 0x00ff) |
34                (((tp->snd_nxt - 1) >> 8) & 0xff00) |
35                (((tp->snd_nxt - 1) & 0xff00) << 8) |
36                (((tp->snd_nxt - 1) & 0x00ff) << 24);

```

### 3.2.4 Detecting Dropped Connections

When a connection is dropped because keepalives were not being acknowledged the TCP callback function will be called with an event code of M\_CLOSED. The socket error will be set to ETIMEDOUT.

```

37  case M_CLOSED:
38      if (so->error == ETIMEDOUT)
39      {

```

```
40      // connection dropped by stack
41  }
42  break;
```

Note that there is no need to call `m_close` as the socket will automatically be closed by the stack.

## 4. UDP

This section contains brief notes on using the UDP layer.

### 4.1 Return Values When Sending UDP Datagrams

The first time `udp_send` is called for a specific host the stack may have to go through the ARP process to find out information about the host. If this is the case then the `udp_send` function will return `ENP_SEND_PENDING`. This is defined with the value of one, and isn't an error condition. Therefore to check for errors returned by the function only negative numbers should be considered.

```
43  If (udp_send(port1, port2, packet) < 0)
44  {
45      // UDP send error
46  }
```

### 4.2 UDP Callback Function

To avoid a memory leak, the `udp_free` function must be called in the UDP callback function when the packet of data has been processed. This is demonstrated in the DHCP client example supplied with the NicheLite stack.

## 5. Sockets

This section contains notes and issues relating to the use of sockets. When a connection is opened or listened for, a socket is created. The socket is a data structure describing the status of the connection and providing a means of transmitting to a remote host.

### 5.1 Freeing Allocated Memory

When a connection is closed the memory allocated to the description of the socket must be freed. It is possible that there is a delay between the connection closing and the memory being freed. If the embedded system is handling enough connection and disconnection requests there could be a memory leak as the list of pending sockets to be freed increases.

It is possible to force the stack to free the memory for the socket immediately. This is handled in the TCP callback function when the event code is `M_CLOSED`. The following code demonstrates how this is achieved.

```
47  case M_CLOSED:
48      if (so->error != ETIMEDOUT)
```

```

49      {
50          so->so_options |= SO_LINGER;
51          so->linger = 0;
52      }
53      break;

```

Note that the SO\_LINGER option should not be used if the stack closed the connection due to a timeout condition.

## 5.2 Runtime Monitoring

The usage of sockets can be monitored during runtime when using a suitable debug interface, such as JTAG or an In-Circuit Emulator.

Access to some internal data is also available via the default serial interface system, however because the interface may affect performance and may not be needed or desired for production firmware, it may be disabled. Using a system such as JTAG provides more flexibility during debugging, and is recommended.

### 5.2.1 Socket Usage

Details of how many sockets are in use and what for may be obtained by watching the msoq linked list defined in tcputil.c.

Name	Value
msoq	struct queue { ... }
q_head	0x40004A10
*q_head	struct q_elt { ... }
qe_next	0x40004B20
q_tail	0x40004C30
q_len	0x00000003
q_max	0x00000003
q_min	0x00000000
<type F2 to edit>	

Fig 3. Watching msoq

The member q\_len is the size of the list, i.e. the number of sockets in use. The q\_max member is the highest number of sockets that have been in use since the stack started executing.

## 6. Stack Configuration

It may be necessary to configure aspects of how the TCP/IP stack behaves for tight integration into an application. This section describes some aspects of this process.

## 6.1 IP Address Configuration

A common requirement for embedded systems that use a TCP/IP stack is to allow for the end user to configure the IP address, and whether the IP address is fixed (static) or obtained from the network (dynamic) using a DHCP server.

For example the system may provide a HTTP or TELNET based interface allowing the IP address to be configured. This is necessary due to the wide variation in networks that the system may be installed into.

The configuration of the IP address is performed in the function `pre_task_setup` in `in_stubs.c`. The following is a code snippet of the relevant section.

```
54  #ifdef DHCP_CLIENT
55  netstatic[i].n_flags |= NF_DHCPC;
56  netstatic[i].n_ipaddr = 0x00000000; /* 0.0.0.0 */
57  #else
58  netstatic[i].n_ipaddr = 0x6400000A; /* 10.0.0.100 */
59  #endif
60  netstatic[i].snmask = 0x000000FF; /* 255.0.0.0 */
61  netstatic[i].n_defgw = 0x0100000A; /* 10.0.0.1 */
62  i++;
```

This section of code configures the `netstatic` array for the Ethernet interface number `i`. If a dynamic IP address is to be used then `NF_DHCPC` is ORed with `n_flags` and the `n_ipaddr` value is ignored. If a static IP address is to be used then it is placed into `n_ipaddr`.

The `snmask` contains the IP network mask and `n_defgw` contains the default gateway for the network.

To allow these values to be configurable these lines need to be replaced with code that can obtain the current configuration, perhaps from non-volatile memory. The following is an example for illustrative purposes.

```
63  // configure static/dynamic ip address based on nvol settings
64  if (NVol_UseStaticIP())
65  {
66      NVol_GetStaticIP(staticip);
67      netstatic[i].n_ipaddr = ((unsigned long)staticip[3] << 24) |
68                             ((unsigned long)staticip[2] << 16) |
69                             ((unsigned long)staticip[1] << 8) |
70                             (unsigned long)staticip[0];
71  }
72  else
73  {
74      netstatic[i].n_flags |= NF_DHCPC;
75      netstatic[i].n_ipaddr = 0x00000000;
76  }
77  i++;
```

## 6.2 MAC Address Configuration

Every Ethernet controller must have a unique MAC address, which is assigned by the IEEE, and this value must be supplied to the TCP/IP stack.

The default hard-coded MAC address for the NicheLite stack is defined in the `eth_info` array, which is initialized in `emac.c`. Because it typically isn't practical to compile different source code for each copy of a product, this method will likely need to be adjusted. For example the `eth_info` array could be located at a specific location in flash memory, allowing patching of the hex file on the production line.

An alternative approach is to store the MAC address in non-volatile memory and supply it to the stack during initialization. This would allow for more flexibility in configuration of the MAC address.

The MAC address should be assigned to the `mac_addr` member of the `eth_info` array inside the `eth_init` function in `emac.c`. The following code snippet is an example for illustrative purposes.

```
78 // get MAC address from NVol module
79 if (!NVol_GetMACAddress((unsigned char *)&(eth_info[dev].mac_addr)))
80 {
81     dprintf("*** failed to get mac address, cannot init ethernet\n");
82     return (-1);
83 }
```

Note that this should be performed before the `mac_addr` member is used for the first time in the `eth_init` function.

## 7. Implementing Raw Send Functionality

Embedded systems often need access to lower level functionality, especially in situations where optimization is required or out-of-the-ordinary operations need to take place.

The NicheLite stack does not provide a single interface for the application to construct IP datagrams from scratch, but it is possible. Note that this approach relies on the internal data structures used by the stack, which could change in future versions.

### 7.1 Example For A Custom Protocol

The following code listing demonstrates how to send an IP datagram containing a custom protocol by directly constructing an IP datagram. The protocol is called MYPROT and contains some data along with an IP-style checksum.

```
1  PACKET p;
2  struct arptabent tp;
3  struct ip *ip_header;
4  char *pprot;
5  struct myprot *myprot_data;
6
7  // allocate memory for packet
8  LOCK_NET_RESOURCE(FREEQ_RESID);
9  p = pk_alloc(ETHHDR_SIZE + sizeof(struct ip) + IP_OPTIONS_LEN + sizeof(struct myprot));
```

```

10  UNLOCK_NET_RESOURCE(FREEQ_RESID);
11  if (!p)
12  {
13      dprintf("Failed to allocate memory for MYPROT packet\n");
14      return;
15  }
16
17  // get start of IP/protocol data buffer
18  p->nb_prot = p->nb_buff + ETHHDR_SIZE;
19  // set length of IP/MYPROT data
20  p->nb_plen = sizeof(struct ip) + IP_OPTIONS_LEN + sizeof(struct myprot);
21  // set destination host
22  p->fhost = MYPROT_HOST;
23  // use first ethernet interface
24  p->net = nets[0];
25
26  // store IP header
27  ip_header = (struct ip *)p->nb_prot;
28  ip_header->ip_ver_ihl = 4 << 4; // IP version 4
29  ip_header->ip_ver_ihl |= (sizeof(struct ip) + IP_OPTIONS_LEN) / 4; // length of IP header in 32-bit
    words
30  ip_header->ip_tos = 0x00; // default
31  ip_header->ip_len = HTONS(24 + sizeof(struct myprot)); // total message length
32  ip_header->ip_id = (unshort)((uid >> 8) | (uid << 8)); // IP datagram ID
33  ip_header->ip_flg_fff = 0x0000; // no flags
34  ip_header->ip_time = 64; // TTL
35  ip_header->ip_prot = MYPROT_PROT; // MYPROT message
36  ip_header->ip_chksum = IPXSUM; // checksum - zero for calculation
37  ip_header->ip_src = m_netp->n_ipaddr; // our IP address
38  ip_header->ip_dest = MYPROT_HOST; // destination
39  // add IP options (IP_OPTIONS_LEN bytes in size)
40  pprot = p->nb_prot + sizeof(struct ip);
41  pprot[0] = 0x94; // router alert option
42  pprot[1] = 0x04;
43  pprot[2] = 0x00;
44  pprot[3] = 0x00;
45  pprot += IP_OPTIONS_LEN;
46
47  // calculate and store IP header checksum
48  ip_header->ip_chksum = ~cksum(ip_header, sizeof(struct ip) + IP_OPTIONS_LEN);
49
50  // increment id for next datagram
51  uid++;
52
53  // store MYPROT data
54  myprot_data = (struct myprot *)pprot;
55  myprot_data->type = 0x01;
56  myprot_data->resptime = 0x02;
57  myprot_data->checksum = IPXSUM;
58
59  // calculate and store MYPROT checksum

```

```
60  myprot_data->checksum = ~cksum(myprot_data, sizeof(struct myprot));
61
62  // destination MAC address
63  tp.t_phy_addr[0] = 0x01;
64  tp.t_phy_addr[1] = 0x00;
65  tp.t_phy_addr[2] = 0x5E;
66  tp.t_phy_addr[3] = 0x7F;
67  tp.t_phy_addr[4] = 0xFF;
68  tp.t_phy_addr[5] = 0xFF;
69
70  // add ethernet header and send it
71  if (et_send(p, &tp) != 0)
72  {
73      dprintf("Failed to send MYPROT message\n");
74      return;
75  }
```

Line 9: allocates memory for the packet. The total size is made up of the Ethernet header, IP header, IP options and MYPROT data.

Lines 18 - 24: some members of the packet need to be initialized. nb\_prot is set to point to the start of the IP header.

Lines 27 – 44: configure the IP header for the message. This includes a TTL value and the protocol number for the custom protocol.

Lines 54 – 57: the data is filled in for the custom protocol.

Lines 63 – 68: the MAC address for the destination is filled in. Alternatively the stack can be instructed to go through the ARP process and automatically determine the destination MAC address.

## 8. Versions Used

This application note was developed with the following software versions:

- NicheLite for LPC revision 1.02
- Keil RealView MDK-ARM version 3.20

## 9. Legal information

### 9.1 Definitions

**Draft** — The document is a draft version only. The content is still under internal review and subject to formal approval, which may result in modifications or additions. NXP Semiconductors does not give any representations or warranties as to the accuracy or completeness of information included herein and shall have no liability for the consequences of use of such information.

### 9.2 Disclaimers

**General** — Information in this document is believed to be accurate and reliable. However, NXP Semiconductors does not give any representations or warranties, expressed or implied, as to the accuracy or completeness of such information and shall have no liability for the consequences of use of such information.

**Right to make changes** — NXP Semiconductors reserves the right to make changes to information published in this document, including without limitation specifications and product descriptions, at any time and without notice. This document supersedes and replaces all information supplied prior to the publication hereof.

**Suitability for use** — NXP Semiconductors products are not designed, authorized or warranted to be suitable for use in medical, military, aircraft, space or life support equipment, nor in applications where failure or malfunction of a NXP Semiconductors product can reasonably be expected to result in personal injury, death or severe property or environmental damage. NXP Semiconductors accepts no liability for inclusion and/or use of NXP Semiconductors products in such equipment or applications and therefore such inclusion and/or use is for the customer's own risk.

**Applications** — Applications that are described herein for any of these products are for illustrative purposes only. NXP Semiconductors makes no representation or warranty that such applications will be suitable for the specified use without further testing or modification.

### 9.3 Trademarks

Notice: All referenced brands, product names, service names and trademarks are property of their respective owners.



10. Contents

1. Introduction .....3

2. Packets.....3

2.1 Packet Allocation and Deallocation .....3

2.1.1 Packet Types and Configuration .....3

2.2 Increasing the Total Number of Packets .....4

2.2.1 Stack Configuration .....4

2.2.2 Testing .....5

2.3 Runtime Monitoring .....5

2.3.1 Detailed Packet Usage.....5

2.3.2 Total Packet Count.....6

2.3.3 Packet Statistics.....6

3. TCP .....7

3.1 Retransmits .....7

3.2 Keepalives.....8

3.2.1 Enabling Keepalives.....8

3.2.2 Adjusting Timing.....8

3.2.3 Fixing Sequence Number Issue .....9

3.2.4 Detecting Dropped Connections .....9

4. UDP .....10

4.1 Return Values When Sending UDP Datagrams .....10

4.2 UDP Callback Function .....10

5. Sockets .....10

5.1 Freeing Allocated Memory .....10

5.2 Runtime Monitoring .....11

5.2.1 Socket Usage.....11

6. Stack Configuration .....11

6.1 IP Address Configuration .....12

6.2 MAC Address Configuration .....13

7. Implementing Raw Send Functionality.....13

7.1 Example For A Custom Protocol .....13

8. Versions Used .....15

9. Legal information .....16

9.1 Definitions .....16

9.2 Disclaimers.....16

9.3 Trademarks.....16

10. Contents.....17

Please be aware that important notices concerning this document and the product(s) described herein, have been included in the section 'Legal information'.



© NXP B.V. 2008. All rights reserved.

For more information, please visit: <http://www.nxp.com>  
For sales office addresses, email to: [salesaddresses@nxp.com](mailto:salesaddresses@nxp.com)